

Final Project: Cache Bypassing Analysis and Performance Study

Brian Park

North Carolina State University, Computer Engineering 786

April 2023

1 Cache Efficiency Analysis

The configuration used is SM2_GTX480. I set the `gpgpu_max_insn` to 100M. And benchmarked from Rodinia are BP, HS, and LUD benchmarks. From ISPASS, I chose BFS, LPS, and NQU.

To determine the category, cache insensitive is classified as any improvement less than 10% or -10%. Otherwise, it is cache unfriendly or cache friendly. The results are shown in Table 1. Since the full description cannot fit in the table, the columns are as follows: benchmark name, kernel name, kernel launch UID, IPC with no cache bypassing, IPC with cache bypassing, percentage change of comparing the IPC with/without bypassing, and the benchmark category.

Table 1: Cache Efficiency Analysis Table

| Benchmark | Kernel | UID | IPC | IPC w/ Bypass | Improvement | Category |
|-----------|---|-----|----------|---------------|-------------|-------------|
| BP | _Z22bpnn_layerforward_CUDAPfS_S_S_ii | 1 | 675.6067 | 671.3728 | -0.63% | Insensitive |
| HS | _Z14calculate_tempiPfS_S_iiiiiiiii | 1 | 701.3718 | 707.6299 | 0.89% | Insensitive |
| LUD | _Z12lud_diagonalPfii | 1 | 0.7026 | 0.7176 | 2.13% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 2 | 9.2446 | 9.1103 | -1.45% | Insensitive |
| LUD | _Z12lud_internalPfii | 3 | 501.2445 | 567.1572 | 13.15% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 4 | 0.7558 | 0.7742 | 2.43% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 5 | 10.9464 | 11.8102 | 7.89% | Insensitive |
| LUD | _Z12lud_internalPfii | 6 | 497.3745 | 574.7466 | 15.56% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 7 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 8 | 10.1697 | 10.9718 | 7.89% | Insensitive |
| LUD | _Z12lud_internalPfii | 9 | 473.0808 | 557.2787 | 17.8% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 10 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 11 | 9.3893 | 10.1287 | 7.87% | Insensitive |
| LUD | _Z12lud_internalPfii | 12 | 462.4784 | 529.6388 | 14.52% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 13 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 14 | 8.6082 | 9.2874 | 7.89% | Insensitive |
| LUD | _Z12lud_internalPfii | 15 | 378.4012 | 504.6895 | 33.37% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 16 | 0.7558 | 0.7742 | 2.43% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 17 | 7.8294 | 8.4467 | 7.88% | Insensitive |
| LUD | _Z12lud_internalPfii | 18 | 357.2093 | 493.7370 | 38.22% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 19 | 0.7558 | 0.7742 | 2.43% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 20 | 7.0473 | 7.6040 | 7.9% | Insensitive |
| LUD | _Z12lud_internalPfii | 21 | 338.0277 | 453.3258 | 34.11% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 22 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 23 | 6.2640 | 6.7609 | 7.93% | Insensitive |
| LUD | _Z12lud_internalPfii | 24 | 324.1251 | 467.1097 | 44.11% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 25 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 26 | 5.4832 | 5.9163 | 7.9% | Insensitive |
| LUD | _Z12lud_internalPfii | 27 | 290.9933 | 405.2074 | 39.25% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 28 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 29 | 4.7006 | 5.0733 | 7.93% | Insensitive |
| LUD | _Z12lud_internalPfii | 30 | 246.8571 | 344.3503 | 39.49% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 31 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 32 | 3.9172 | 4.2288 | 7.95% | Insensitive |
| LUD | _Z12lud_internalPfii | 33 | 208.6225 | 253.2766 | 21.4% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 34 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 35 | 3.1348 | 3.3833 | 7.93% | Insensitive |
| LUD | _Z12lud_internalPfii | 36 | 142.2966 | 172.1319 | 20.97% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 37 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 38 | 2.3514 | 2.5387 | 7.97% | Insensitive |
| LUD | _Z12lud_internalPfii | 39 | 111.9498 | 134.8471 | 20.45% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 40 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 41 | 1.5679 | 1.6926 | 7.95% | Insensitive |
| LUD | _Z12lud_internalPfii | 42 | 39.4499 | 44.9208 | 13.87% | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 43 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| LUD | _Z13lud_perimeterPfii | 44 | 0.8583 | 0.8467 | -1.35% | Insensitive |
| LUD | _Z12lud_internalPfii | 45 | 16.2623 | 16.6957 | 2.67% | Insensitive |
| LUD | _Z12lud_diagonalPfii | 46 | 0.7558 | 0.7741 | 2.42% | Insensitive |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 1 | 217.5687 | 167.9066 | -22.83% | Friendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 2 | 206.0139 | 146.9099 | -28.69% | Friendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 3 | 165.9271 | 112.0179 | -32.49% | Friendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 4 | 76.2236 | 61.3361 | -19.53% | Friendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 5 | 21.3021 | 36.1667 | 69.78% | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 6 | 22.5533 | 44.4395 | 97.04% | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 7 | 46.5675 | 86.5094 | 85.77% | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 8 | 354.4445 | 455.3303 | 28.46% | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 9 | 473.1056 | 486.7920 | 2.89% | Insensitive |
| LPS | _Z13GPU_laplace3diiiPfS_ | 1 | 383.1095 | 408.8568 | 6.72% | Insensitive |
| NQU | _Z24solve_nqueen_cuda_kerneliiPjS_S_S_i | 1 | 30.4185 | 30.7699 | 1.16% | Insensitive |

2 Profiling-Based Cache Bypassing

I implemented what was mentioned in the bypassing paper [1]. Surprisingly, the implementation was very simple. I add two data structures, one for storing SM, kernel UID, address information and the reference count in a nested C++ map. I create two of these nested tables, one that is recorded during profiled run and dumped into `profile_dump.out` and one that is reused in the second or future runs where it reads data from `profile_dump.out`. This is simply added in `ldst_unit::memory_cycle` where memory requests are handled. If a memory request has a reference counter less than 3, I bypass the L1D cache. This turns out to be a simple 3 line addition that profiles and does the cache bypassing if a profiled run is available. Additional logic needs to be added in `gpgpusim_entrypoint.cc` where we have to log to `profile_dump.out` and read from `profile_dump.out` using some file and string manipulations. For `git diff` of the implementation, it is available at the bottom of this report.

The benchmarks that I reran and reanalyzed are the ones that were **cache unfriendly**. Those included kernels in LUD and BFS. Table 2 shows the IPC with profiling from V100 configuration. Note that it is not exactly fair to compare against SM2 GTX480, but as the spec stated, SM2 configuration with profiling caused some deadlocks. But, the results look sound and we see that there are some improvements and degradation in IPC with profiling.

Table 2: Cache Efficiency Analysis Table with Profiling

| Benchmark | Kernel | UID | IPC | IPC w/ Bypass | IPC w/ Profile | Category |
|-----------|-------------------------------|-----|----------|---------------|----------------|-------------|
| LUD | _Z12lud_diagonalPfii | 1 | 0.7026 | 0.7176 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 2 | 9.2446 | 9.1103 | 15.3590 | Insensitive |
| LUD | _Z12lud_internalPfii | 3 | 501.2445 | 567.1572 | 721.2603 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 4 | 0.7558 | 0.7742 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 5 | 10.9464 | 11.8102 | 14.3418 | Insensitive |
| LUD | _Z12lud_internalPfii | 6 | 497.3745 | 574.7466 | 642.5734 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 7 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 8 | 10.1697 | 10.9718 | 13.3236 | Insensitive |
| LUD | _Z12lud_internalPfii | 9 | 473.0808 | 557.2787 | 557.6649 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 10 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 11 | 9.3893 | 10.1287 | 12.3045 | Insensitive |
| LUD | _Z12lud_internalPfii | 12 | 462.4784 | 529.6388 | 506.7778 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 13 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 14 | 8.6082 | 9.2874 | 11.3033 | Insensitive |
| LUD | _Z12lud_internalPfii | 15 | 378.4012 | 504.6895 | 429.6447 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 16 | 0.7558 | 0.7742 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 17 | 7.8294 | 8.4467 | 10.2805 | Insensitive |
| LUD | _Z12lud_internalPfii | 18 | 357.2093 | 493.7370 | 357.2629 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 19 | 0.7558 | 0.7742 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 20 | 7.0473 | 7.6040 | 9.2568 | Insensitive |
| LUD | _Z12lud_internalPfii | 21 | 338.0277 | 453.3258 | 293.3001 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 22 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 23 | 6.2640 | 6.7609 | 8.2322 | Insensitive |
| LUD | _Z12lud_internalPfii | 24 | 324.1251 | 467.1097 | 235.2133 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 25 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 26 | 5.4832 | 5.9163 | 7.2186 | Insensitive |
| LUD | _Z12lud_internalPfii | 27 | 290.9933 | 405.2074 | 180.4473 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 28 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 29 | 4.7006 | 5.0733 | 6.1903 | Insensitive |
| LUD | _Z12lud_internalPfii | 30 | 246.8571 | 344.3503 | 132.7995 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 31 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 32 | 3.9172 | 4.2288 | 5.1610 | Insensitive |
| LUD | _Z12lud_internalPfii | 33 | 208.6225 | 253.2766 | 92.7536 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 34 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 35 | 3.1348 | 3.3833 | 4.1308 | Insensitive |
| LUD | _Z12lud_internalPfii | 36 | 142.2966 | 172.1319 | 59.3623 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 37 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 38 | 2.3514 | 2.5387 | 3.1034 | Insensitive |
| LUD | _Z12lud_internalPfii | 39 | 111.9498 | 134.8471 | 33.5850 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 40 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 41 | 1.5679 | 1.6926 | 2.0726 | Insensitive |
| LUD | _Z12lud_internalPfii | 42 | 39.4499 | 44.9208 | 14.9430 | Unfriendly |
| LUD | _Z12lud_diagonalPfii | 43 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| LUD | _Z13lud_perimeterPfii | 44 | 0.8583 | 0.8467 | 1.0381 | Insensitive |
| LUD | _Z12lud_internalPfii | 45 | 16.2623 | 16.6957 | 3.7422 | Insensitive |
| LUD | _Z12lud_diagonalPfii | 46 | 0.7558 | 0.7741 | 0.7782 | Insensitive |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 1 | 217.5687 | 167.9066 | 104.4680 | Friendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 2 | 206.0139 | 146.9099 | 90.2066 | Friendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 3 | 165.9271 | 112.0179 | 77.9455 | Friendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 4 | 76.2236 | 61.3361 | 72.6998 | Friendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 5 | 21.3021 | 36.1667 | 66.7220 | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 6 | 22.5533 | 44.4395 | 66.8945 | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 7 | 46.5675 | 86.5094 | 111.0300 | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 8 | 354.4445 | 455.3303 | 175.9334 | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 9 | 473.1056 | 486.7920 | 149.3736 | Insensitive |

To see the improvements on cache unfriendly kernels better, Table 3 is shown with only the cache unfriendly data filtered out. We see that 6 out of 18 kernels that are cache unfriendly are improved in IPC with profiling. It may seem low, but if we look at some of the cache insensitive kernels, many of them are noticeably improved. It seems that

the 10% limit set by me could make these results look skewed, but it also tells us that in the extreme case of cache unfriendly kernels with 10% difference with bypassing/no bypassing may not bring drastic improvements.

Table 3: Cache Efficiency Analysis Table with Profiling (Cache Unfriendly Kernels Only)

| Benchmark | Kernel | UID | IPC | IPC w/ Bypass | IPC w/ Profile | Category |
|-----------|-------------------------------|-----|----------|---------------|----------------|------------|
| LUD | _Z12lud_internalPfii | 3 | 501.2445 | 567.1572 | 721.2603 | Unfriendly |
| LUD | _Z12lud_internalPfii | 6 | 497.3745 | 574.7466 | 642.5734 | Unfriendly |
| LUD | _Z12lud_internalPfii | 9 | 473.0808 | 557.2787 | 557.6649 | Unfriendly |
| LUD | _Z12lud_internalPfii | 12 | 462.4784 | 529.6388 | 506.7778 | Unfriendly |
| LUD | _Z12lud_internalPfii | 15 | 378.4012 | 504.6895 | 429.6447 | Unfriendly |
| LUD | _Z12lud_internalPfii | 18 | 357.2093 | 493.7370 | 357.2629 | Unfriendly |
| LUD | _Z12lud_internalPfii | 21 | 338.0277 | 453.3258 | 293.3001 | Unfriendly |
| LUD | _Z12lud_internalPfii | 24 | 324.1251 | 467.1097 | 235.2133 | Unfriendly |
| LUD | _Z12lud_internalPfii | 27 | 290.9933 | 405.2074 | 180.4473 | Unfriendly |
| LUD | _Z12lud_internalPfii | 30 | 246.8571 | 344.3503 | 132.7995 | Unfriendly |
| LUD | _Z12lud_internalPfii | 33 | 208.6225 | 253.2766 | 92.7536 | Unfriendly |
| LUD | _Z12lud_internalPfii | 36 | 142.2966 | 172.1319 | 59.3623 | Unfriendly |
| LUD | _Z12lud_internalPfii | 39 | 111.9498 | 134.8471 | 33.5850 | Unfriendly |
| LUD | _Z12lud_internalPfii | 42 | 39.4499 | 44.9208 | 14.9430 | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 5 | 21.3021 | 36.1667 | 66.7220 | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 6 | 22.5533 | 44.4395 | 66.8945 | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 7 | 46.5675 | 86.5094 | 111.0300 | Unfriendly |
| BFS | _Z6KernelP4NodePiPbS2_S1_S2_i | 8 | 354.4445 | 455.3303 | 175.9334 | Unfriendly |

To reproduce the results from the analysis, all you have to do is run the same benchmarks twice, but copying the new V100 config over. For convenience, a bash script is provided, but it still requires to be run twice in order to produce the proper results.

Please make sure the root directory has this structure with the script:

```
ls ~
ISPASS
analysis.sh
env
gpgpu-sim_distribution
rodinia
```

Then you can run `analysis.sh` twice, making sure to edit the configs and changes to GPGPU-Sim properly:

```
bash analysis.sh
bash analysis.sh

diff --git a/src/gpgpu-sim/shader.cc b/src/gpgpu-sim/shader.cc
index c6e7b8f..da102a5 100644
--- a/src/gpgpu-sim/shader.cc
+++ b/src/gpgpu-sim/shader.cc
@@ -52,6 +52,8 @@
#define MAX(a, b) (((a) > (b)) ? (a) : (b))
#define MIN(a, b) (((a) < (b)) ? (a) : (b))

+std::map<unsigned, std::map<unsigned, std::map<unsigned, unsigned>>> SM_kernel_addr_table;
+
mem_fetch *shader_core_mem_fetch_allocator::alloc(
    new_addr_type addr, mem_access_type type, unsigned size, bool wr,
    unsigned long long cycle) const {
@@ -2039,6 +2041,12 @@ bool ldst_unit::memory_cycle(warp_inst_t &inst,
    const mem_access_t &access = inst.accessq_back();

    bool bypassL1D = false;
+ SM_kernel_addr_table[m_core->get_sid()][m_core->get_kernel()->get_uid()][access.get_addr()]++;
```

```

+
+ if (profile_table[m_core->get_sid()][m_core->get_kernel()->get_uid()][access.get_addr()] < 3){
+     bypassL1D = true;
+ }
+
+ if (CACHE_GLOBAL == inst.cache_op || (m_L1D == NULL)) {
+     bypassL1D = true;
+ } else if (inst.space.is_global()) { // global memory access
diff --git a/src/gpgpu-sim/shader.h b/src/gpgpu-sim/shader.h
index 6481790..ed34d41 100644
--- a/src/gpgpu-sim/shader.h
+++ b/src/gpgpu-sim/shader.h
@@ -70,6 +70,9 @@

#define WRITE_MASK_SIZE 8

+extern std::map<unsigned, std::map<unsigned, std::map<unsigned, unsigned>>> SM_kernel_addr_table;
+extern std::map<unsigned, std::map<unsigned, std::map<unsigned, unsigned>>> profile_table;
+
+ class gpgpu_context;

+ enum exec_unit_type_t {
diff --git a/src/gpgpusim_entrystatepoint.cc b/src/gpgpusim_entrystatepoint.cc
index f4287d8..9610b0b 100644
--- a/src/gpgpusim_entrystatepoint.cc
+++ b/src/gpgpusim_entrystatepoint.cc
@@ -28,6 +28,10 @@

#include "gpgpusim_entrystatepoint.h"
#include <stdio.h>
+#include <iostream>
+#include <fstream>
+#include <string>

#include "../libcuda/gpgpu_context.h"
#include "cuda-sim/cuda-sim.h"
@@ -42,6 +42,7 @@

static int sg_argc = 3;
static const char *sg_argv[] = {"", "-config", "gpgpusim.config"};
+std::map<unsigned, std::map<unsigned, std::map<unsigned, unsigned>>> profile_table;

void *gpgpu_sim_thread_sequential(void *ctx_ptr) {
    gpgpu_context *ctx = (gpgpu_context *)ctx_ptr;
@@ -229,6 +229,20 @@
    sem_init(&(the_gpgpusim->g_sim_signal_finish), 0, 0);
    sem_init(&(the_gpgpusim->g_sim_signal_exit), 0, 0);

+ std::ifstream profile_log;
+ profile_log.open("profile_dump.out");
+ if(profile_log.is_open()){
+     std::string line;
+     size_t SM_id, kernel_uid, addr, ref_cnt;
+     while(getline(profile_log, line)){
+         if(line.empty())
+             continue;
+         sscanf(line.c_str(), "SM: %lu, kernel: %lu, addr: %lu, ref: %lu", &SM_id, &kernel_uid, &addr, &ref_cnt);

```

```

+     profile_table[SM_id][kernel_uid][addr] = ref_cnt;
+ }
+ profile_log.close();
+ }
+
+ return the_gpgpusim->g_the_gpu;
+ }

@@ -266,6 +285,18 @@ void gpgpu_context::print_simulation_time() {
    printf("gpgpu_simulation_rate = %u (cycle/sec)\n", cycles_per_sec);
    printf("gpgpu_silicon_slowdown = %ux\n",
           the_gpgpusim->g_the_gpu->shader_clock() * 1000 / cycles_per_sec);
+
+ std::ofstream profile_log;
+ profile_log.open("profile_dump.out");
+ for (auto SM = SM_kernel_addr_table.begin(); SM != SM_kernel_addr_table.end(); SM++) {
+     for (auto kernel = SM->second.begin(); kernel != SM->second.end(); kernel++) {
+         for (auto addr = kernel->second.begin(); addr != kernel->second.end(); addr++) {
+             profile_log << "SM: " << SM->first << ", kernel: " << kernel->first << ", addr: " << addr->first <
+         }
+     }
+ }
+ profile_log.close();
+
+ fflush(stdout);
+ }

```

References

- [1] C. Li, S. L. Song, H. Dai, A. Sidelnik, S. K. S. Hari, and H. Zhou, “Locality-driven dynamic gpu cache bypassing,” in *Proceedings of the 29th ACM on International Conference on Supercomputing, ICS ’15*, (New York, NY, USA), p. 67–77, Association for Computing Machinery, 2015.